

## ELK:

### Elastic search

- Elasticsearch is a NoSQL search engine.
- It provides a distributed, full-text search engine with an HTTP web interface and of JSON documents.
- Elasticsearch is developed in Java and is released as open source under the terms of the Apache License.
- Official clients are available in Java, .NET (C#), PHP, Python, Apache Groovy, Ruby and many other languages.
- According to the DB-Engines ranking, Elasticsearch is the most popular enterprise search engine.
- Elasticsearch does not have tables, and a schema is not required.
- Elasticsearch stores data documents that consist of JSON strings inside an index.
- Elasticsearch is a distributed, **RESTful** search and analytics engine.
- Elasticsearch is a NRT (near real time) search platform.
- What this means is there is a slight latency (normally one second) from the time you index a document until the time it becomes searchable.

## Terminology

- **Cluster**

- A cluster is a collection of one or more nodes (servers) that together holds your entire data and provides federated indexing and search capabilities across all nodes.

- **Node**

- A node is a single server that is part of your cluster, stores your data, and participates in the cluster's indexing and search capabilities.

- **Index**

- An index is a collection of documents that have somewhat similar characteristics.
- For example, you can have an index for customer data, another index for a product catalog, and yet another index for order data.
- An index is identified by a name (that must be all lowercase) and this name is used to refer to the index when performing indexing, search, update, and delete operations against the documents in it.
- In a single cluster, you can define as many indexes as you want.

- **Document**

- A document is a basic unit of information that can be indexed.
- For example, you can have a document for a single customer, another document for a single product, and yet another for a single order.
- This document is expressed in JSON (JavaScript Object Notation).


- **Shards and replicas**

- Every time you index a document elasticsearch will decide which primary shard is supposed to hold that document and will index it there.
- Having multiple shards helps taking advantage of parallel processing on a single machine, but the whole point is that if we start another elasticsearch instance on the same cluster, the shards will be distributed in an even way over the cluster.
- Every elasticsearch index is composed of at least one primary shard since that's where the data is stored.
- Another type of shard is a replica. The default is 1, meaning that every primary shard will be copied to another shard that will contain the same data.
- Replicas are used to increase search performance and for fail-over.

## Setup:

- Elastic download is available in:  
<https://www.elastic.co/downloads/elasticsearch>

- Download the zip file



- Once download is done, extract the zip file to a known location
- Open the bin folder (e.g. “cd <Elastic search unzipped folder>\bin”) and run **elasticsearch** file (or **elasticsearch.bat** on Windows) and wait until you see: **started**
- Open any browser and navigate to your local elastic search engine at <http://localhost:9200/>
- The engine information will show.
- Currently we have no data stored.

## CRUD operations (create, read, update, delete):

- **Create**

- To create a new Document, we use an POST request method.
- We can post to the following server address:  
[http://localhost:9200/my\\_app/users/1](http://localhost:9200/my_app/users/1)
- The path can be anything, but we will need to remember it for getting this data later..

- **Read**

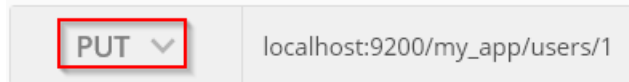
- To query data we can go to a specific index as we did earlier:  
**localhost:9200/my\_app/users/1**
- Another way to retrieve all relevant data will be using the **\_search** which will return all available data from a specific point.
- For example the following URL will return all users:  
**localhost:9200/my\_app/users/\_search**
- We can also use query parameters using “?q=” syntax in Query to filter the records:  
**localhost:9200/my\_app/users/\_search?q=\_id:2**

**OR**

**localhost:9200/my\_app/users/\_search?q=name:John**

- **Update**

- To update an existing document data, we will use **PUT** method:
- The address will be the same as the “posting address” (localhost:9200/my\_app/users/1) except this time we will be using **PUT** method:

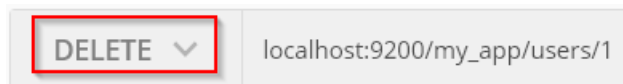


A screenshot of a REST client interface. On the left, there is a dropdown menu with 'PUT' selected and a small downward arrow. This dropdown is highlighted with a red rectangular border. To the right of the dropdown, the URL 'localhost:9200/my\_app/users/1' is displayed in a light gray box.

- Inside the body, simply change anything you want to update and press **send**
- To simply test result, run search again (localhost:9200/my\_app/users/\_search)

- **Delete**

- To delete an existing document data, we will use **DELETE** method:
- The address will be the same as the “posting address” (localhost:9200/my\_app/users/1) except this time we will be using **DELETE** method:



A screenshot of a REST client interface. On the left, there is a dropdown menu with 'DELETE' selected and a small downward arrow. This dropdown is highlighted with a red rectangular border. To the right of the dropdown, the URL 'localhost:9200/my\_app/users/1' is displayed in a light gray box.

- In delete, there no reason to add request body..
- In the example above we changed a user name
- To simply test result, run search again (localhost:9200/my\_app/users/\_search)

## **Kibana**

- Kibana is an open source data visualization plugin for Elasticsearch.
- It provides visualization capabilities on top of the content indexed on an Elasticsearch cluster.
- Users can create bar, line and scatter plots, or pie charts and maps on top of large volumes of data.
- The combination of Elasticsearch, Logstash, and Kibana, referred to as the "Elastic Stack" (formerly the "ELK stack"), is available as a product or service.
- Logstash provides an input stream to Elastic for storage and search, and Kibana accesses the data for visualizations such as dashboards.

## Setup:

- Kibana download is available in:  
<https://www.elastic.co/downloads/kibana>
- Choose your OS

Downloads: [⬆ WINDOWS sha](#) [⬆ MAC sha](#) [⬆ LINUX 64-BIT sha](#)  
[⬆ RPM 64-BIT sha](#) [⬆ DEB 64-BIT sha](#)

- Once download is done, extract the zip file to a known location
- Open the config folder (e.g. "cd <Kibana unzipped folder>\config") and open **config/kibana.yml** in notepad/++
- Find the line contains **elasticsearch.url** and remove the hash(#) from the beginning of the line and save the file.

```
# The URL of the Elasticsearch instance to use for all your queries.  
#elasticsearch.url: "http://localhost:9200"
```

- This will point Kibana to our Elastic search instance.
- Now go to bin folder and run **bin/kibana** (or **bin\kibana.bat** on Windows)
- Open any browser and navigate to your local elastic search engine at <http://localhost:5601/>

```
[info][listening] Server running at http://localhost:5601
```





- Once you are inside Kibana management tab, you will need to define what the index pattern is, you want Kibana to “listen” to.
  - Simply use the one from earlier like the following: my\_app/\*
  - You should then see success message
  - You can now press **next step** and then **create index pattern** and wait until indexing is done.
  - Now, we can go to visualize tab → Create new visualization for the example we will choose Metric.
  - Choose your index and give it a name, and press save button
  - To see the results immediately, we can go to Postman and create more records and then press refresh.
- \*\* We can create as many dashboards / visualization as we need

## LogStash

- Logstash is an open source, server-side data processing pipeline that ingests data from a multitude of sources simultaneously, transforms it, and then sends it to your favorite “stash.” (Ours is Elasticsearch, naturally.)
- Logstash is used to gather logging messages, convert them into json documents and store them in an ElasticSearch cluster.
- Of course those can be viewed in Kibana, which makes the ELK stack so powerful

## Setup:

- Logstash download is available in:  
<https://www.elastic.co/downloads/logstash>
- Download the zip file

Downloads:  ZIP sha  TAR sha  DEB sha

- Once download is done, extract the zip file to a known location
- Open the bin folder (e.g. “cd <Logstash unzipped folder>\bin”) and create a file **logstash.conf** the following configuration should be inside:

```
input {
  file {
    path => "C:/LogStash/myLogs/log_file.txt"
  }
}
output {
  stdout {
    codec => rubydebug
  }
  elasticsearch {
    hosts => ["localhost:9200"]
    action => "index"
    index => "log"
  }
}
```

Logs input source

Using rubydebug to see logs in console

Elastic search path and index for the data to be stored

## Grok

- **Grok** is a Logstash plugin used to parse arbitrary text and structure it.
- **Grok** is a great way to parse unstructured log data into something structured and queryable.
- This tool is perfect for syslog logs, apache and other webserver logs, mysql logs, and in general, any log format that is generally written for humans and not computer consumption.
- Logstash ships with about 120 patterns by default. You can find them here: <https://github.com/logstash-plugins/logstash-patterns-core/tree/master/patterns>
- Let's take an easy server access log which holds the IP address and the method of the request:

**192.168.99.100 GET**

- To filter this type of messages, we can add a filter using grok inside **Logstash.conf** which will look like that:

```
filter {  
  grok {  
    match => { "message" => "%{IP:client} %{WORD:method}" }  
  }  
}
```

---

## Grafana

- Grafana is an open-source, general purpose dashboard and graph composer, which runs as a web application.
- Grafana allows you to query, visualize, alert on and understand your metrics no matter where they are stored.
- Create, explore, and share dashboards with your team and foster a data driven culture.
- Our data source will be elastic search

## Setup:

- Download **Grafana** at:  
<http://docs.grafana.org/installation/windows/>
- Extract the zip to a known location
- Double click **grafana-server.exe**
- Enter **Grafana** default address at **localhost:3000**
- Default user name and password are **admin**
- Once entered, you will be asked to change password it for security reasons.

- Connect data source by pressing add data source
- Then, add the below configurations.
- We will use **log** index name to see the Logstash index we created earlier.
- Press save and test button.
- If everything is OK, you should see “Index OK”

The screenshot shows the Grafana configuration page for an Elasticsearch data source. The interface is dark-themed. On the left, under the 'Name' field, 'my\_data' is entered. The 'Type' is set to 'Elasticsearch'. Under the 'HTTP' section, the 'URL' is 'http://localhost:9200'. On the right, under 'Elasticsearch details', the 'Index name' is 'log', the 'Time field name' is '@timestamp', and the 'Version' is '5.6+'. The 'Max concurrent shard requests' is set to 256, and the 'Min time interval' is 10s.

- Like Kibana, Grafana is a visualization tool, so we will need to setup our dashboard.
- To create a new dashboard/s, press **Dashboard → Home → New dashboard → Graph → Panel options → Edit and add a new graph.**
- Grafana has an alert system, so you can configure alerts to be sent in real time
- To configure the alert system go to **/conf** folder and open **defaults.ini** file to edit.
- Once email is setup, go to test it by adding a recipient here:
- This will only work if "**Less secured apps**" is configured in your Gmail account: <https://myaccount.google.com/u/1/security>